

Machine Learning & Artificial Intelligence for Data Scientists: Classification (Part2)

Ke Yuan

<https://kyuanlab.org/>

School of Computing Science

Logistic Regression

- — — ▶ Alternative is to directly model $P(T_{\text{new}} = k | \mathbf{x}_{\text{new}}, \mathbf{X}, \mathbf{t}) = f(\mathbf{x}_{\text{new}}; \mathbf{w})$ with some parameters \mathbf{w} .
- ▶ We've seen $f(\mathbf{x}_{\text{new}}; \mathbf{w}) = \mathbf{w}^T \mathbf{x}_{\text{new}}$ before – can we use it here?
 - ▶ No – *output is unbounded and so can't be a probability.*
- ▶ But, can use $P(T_{\text{new}} = k | \mathbf{x}_{\text{new}}, \mathbf{w}) = h(f(\mathbf{x}_{\text{new}}; \mathbf{w}))$ where $h(\cdot)$ *squashes* $f(\mathbf{x}_{\text{new}}; \mathbf{w})$ to lie between 0 and 1 – a probability.

Recap on probability

— — —

- ▶ Discrete v continuous.
- ▶ Probabilities and densities.
- ▶ Joint probabilities and densities.
- ▶ Independence.
- ▶ Conditioning.

Random variables

If I toss a coin and assign the variable X the value 1 if the coin lands heads and 0 if it lands tails, X is a random variable.

We don't know which value X will take but we do know the possible values and how likely they are.

Discrete and continuous RVs

- - ▶ Random events with outcomes that we can count:
Discrete.
 - ▶ Coin toss.
 - ▶ Rolling a die.
 - ▶ Next word in a document.
 - ▶ Number of emails sent in a day.
 - ▶ Random events with outcomes that we cannot count
Continuous.
 - ▶ Winning time in Olympic 100m.

Discrete and continuous RVs

Definitions

Random variables given capital letters - X , Y .

Lower case letters used for values they can take - x , y .

Discrete RVs

— — —

Discrete RVs defines by probabilities of different events taking place. E.g. probability of random variable X taking value x :

$$P(X = x)$$

For example, fair coin:

$$P(X = 1) = 0.5, P(X = 0) = 0.5$$

Die:


$$P(Y = y) = 1/6$$

Probabilities are constrained:

$$0 \leq P(Y = y) \leq 1, \sum_y P(Y = y) = 1.$$

Continuous RVs

- ▶ Don't define probabilities of particular outcomes as we can't count them!
- ▶ Instead define a density function $p(x)$:



$p(x)$ tells us how likely different values are. These are **not** probabilities!

- ▶ We can compute probabilities of ranges by computing the area under the curve:

$$P(6 \leq X \leq 8) = \int_{x=6}^{x=8} p(x) dx$$

- ▶ Densities are constrained:

$$p(x) \geq 0, \quad \int_{-\infty}^{\infty} p(x) dx = 1$$


Joint probabilities and densities

Joint probabilities

For two discrete RVs, X and Y , $P(X = x, Y = y)$ is the probability that RV X has value x **and** RV Y has value y .

Joint densities

For two continuous RVs, x_0 and x_1 , $p(x_0, x_1)$ is the joint density:



Dependence/Independence

— — —

- ▶ Let X be the random variable for the toss of a coin (1=heads, 0=tails)
- ▶ Let Y be the random variable for the rolling of a die.
- ▶ $P(X = 1, Y = 3)$ is the probability that I will roll a head **and** a 3.
- ▶ The outcome of X does not depend on Y .
- ▶ X and Y are independent.

$$P(X = x, Y = y) = P(X = x)P(Y = y)$$

Dependence/Independence

— — —

- ▶ Let X be the random variable for the event – I'm playing tennis (1=yes, 0=no)
- ▶ Let Y be the random variable for the event – It is raining (1=yes, 0=no)
- ▶ $P(X = 1, Y = 1)$ is the probability that I am playing and it is raining.
- ▶ The outcome of X **does** depend on Y .
- ▶ X and Y are dependent.

$$P(X = x, Y = y) \neq P(X = x)P(Y = y)$$

Conditioning

— — —

- ▶ Let X be the random variable for the event – I'm playing tennis (1=yes, 0=no)
- ▶ Let Y be the random variable for the event – It is raining (1=yes, 0=no)
- ▶ Because they are dependent, we can work with **conditional** probabilities.
- ▶ e.g. the probability that I am playing **given that** it is raining:

$$P(X = 1|Y = 1)$$

- ▶ Allows us to decompose the joint probability:

$$P(X = x, Y = y) = P(X = x|Y = y)P(Y = y)$$

Conditioning - continuous

— — —

Example 1:

$$p(t_n | x_n, \mathbf{w})$$

This is the density of t_n conditioned on a particular value of x and our model parameters \mathbf{w} .

Example 2:

$$P(9 \leq t_n \leq 10 | x_n, \mathbf{w})$$

This is the probability of t_n being between 9 and 10 conditioned on a particular value of x and our model parameters \mathbf{w} .


Notational nuance

Technically, we should write $p(t_n | X = x_n, W = \mathbf{w})$ but this becomes unwieldy and gets confusing (difference between X and \mathbf{X} ?). So, we'll use $p(t_n | x_n, \mathbf{w})$.

- — —
- For logistic regression (binary), we use the sigmoid function:

$$P(T_{\text{new}} = 1 | \mathbf{x}_{\text{new}}, \mathbf{w}) = h(\mathbf{w}^T \mathbf{x}_{\text{new}}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x}_{\text{new}})}$$

**Back to logistic
regression:
Sigmoid function**



Introducing Likelihood of a single label

$$\text{if } t_n = 1, \quad p(t_n = 1 | \mathbf{x}_n, \mathbf{w}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x}_n)}$$

$$\text{if } t_n = 0, \quad p(t_n = 0 | \mathbf{x}_n, \mathbf{w}) = 1 - p(t_n = 1 | \mathbf{x}_n, \mathbf{w})$$

One formula for both scenarios

— — —

Likelihood function for nth data point

$$p(t_n | \mathbf{x}_n, \mathbf{w}) = p(t_n = 1 | \mathbf{x}_n, \mathbf{w})^{t_n} (1 - p(t_n = 1 | \mathbf{x}_n, \mathbf{w}))^{(1-t_n)}$$

Likelihood function for all data points

- Assuming data points are independent of each other

$$\text{Likelihood: } p(\mathbf{t}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(t_n|\mathbf{x}_n, \mathbf{w})$$

$$\text{Log-Likelihood: } \log p(\mathbf{t}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N \log p(t_n|\mathbf{x}_n, \mathbf{w})$$

What about Loss?

$$\text{Log-Likelihood}(\mathbf{t}, \mathbf{X}; \mathbf{w}) = \sum_{n=1}^N \log p(t_n | \mathbf{x}_n, \mathbf{w})$$

$$\text{Loss}(\mathbf{t}, \mathbf{X}; \mathbf{w}) = -\text{Log-Likelihood}(\mathbf{t}, \mathbf{X}; \mathbf{w}) = - \sum_{n=1}^N \log p(t_n | \mathbf{x}_n, \mathbf{w})$$

Find the optimal parameters


$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} \operatorname{Loss}(\mathbf{t}, \mathbf{X}; \mathbf{w})$$

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmax}} \operatorname{Likelihood}(\mathbf{t}, \mathbf{X}; \mathbf{w})$$

$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmax}} \operatorname{Log-Likelihood}(\mathbf{t}, \mathbf{X}; \mathbf{w})$$

```
In [59]: from sklearn.linear_model import LogisticRegression
clf = LogisticRegression().fit(X, t)
Z = clf.predict_proba(np.c_[xx.ravel(), yy.ravel()]][:, 1]
# Put the result into a color plot
Z = Z.reshape(xx.shape)
plt.pcolormesh(xx, yy, Z, cmap=cmap_light)
# Plot also the training points
plt.scatter(X[:, 0], X[:, 1], c=t, cmap=cmap_bold,
            edgecolor='k', s=100)
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())
plt.title("Logistic Regression")
mean_cv_score = np.mean( cross_val_score(clf, X, t, cv=5) )
print("5-fold average CV error:", 1-mean_cv_score)
```

5-fold average CV error: 0.0250000000000000022



**Example on orange and
lemon data**

Can be regularised just the same as linear regression


$$\hat{\mathbf{w}}_{l2} = \underset{\mathbf{w}}{\operatorname{argmin}} \operatorname{Loss}(\mathbf{t}, \mathbf{X}; \mathbf{w}) + \frac{1}{C} \mathbf{w}^T \mathbf{w}$$

$$\hat{\mathbf{w}}_{l1} = \underset{\mathbf{w}}{\operatorname{argmin}} \operatorname{Loss}(\mathbf{t}, \mathbf{X}; \mathbf{w}) + \frac{1}{C} \sum_d |w_d|$$

```
In [53]: parameters = {'C':cs}
logit_reg = LogisticRegression(penalty='l2', tol=1e-5, max_iter=1e4)
clf = GridSearchCV(logit_reg, parameters, cv=5)
clf.fit(X,t)

Z = clf.predict_proba(np.c_[xx.ravel(), yy.ravel()])[:, 1]
Z = Z.reshape(xx.shape) # Put the result into a color plot
plt.pcolormesh(xx, yy, Z, cmap=cmap_light)
plt.scatter(X[:, 0], X[:, 1], c=t, cmap=cmap_bold, edgecolor='k', s=100) # Plot all
so the training points
plt.xlim(xx.min(), xx.max())
plt.ylim(yy.min(), yy.max())
plt.title("Logistic Regression with l2 regularisation")
print("5-fold average CV error:", 1-clf.best_score_)
```

5-fold average CV error: 0.0250000000000000022



```
In [49]: from sklearn.svm import l1_min_c
from sklearn.model_selection import GridSearchCV
cs = l1_min_c(X, t, loss='log')*np.logspace(0, 7, 16)
print(cs)

[2.89435601e-02  8.47653998e-02  2.48247728e-01  7.27029358e-01
 2.12921058e+00  6.23570098e+00  1.82621518e+01  5.34833516e+01
 1.56633727e+02  4.58724513e+02  1.34344105e+03  3.93446133e+03
 1.15226388e+04  3.37457135e+04  9.88292004e+04  2.89435601e+05]
```

L2 regularised Logistic Regression

Performance Evaluations

— — —

- ▶ We've seen 2 classification algorithms.
- ▶ How do we choose?
 - ▶ Which algorithm?
 - ▶ Which parameters?
- ▶ Need performance indicators.
- ▶ We'll cover:
 - ▶ 0/1 loss.
 - ▶ ROC analysis (sensitivity and specificity)
 - ▶ Confusion matrices

0/1 loss

— — —

- ▶ 0/1 loss: proportion of times classifier is wrong.
- ▶ Consider a set of predictions t_1, \dots, t_N and a set of true labels t_1^*, \dots, t_N^* .
- ▶ Mean loss is defined as:

$$\frac{1}{N} \sum_{n=1}^N \delta(t_n \neq t_n^*)$$

- ▶ ($\delta(a)$ is 1 if a is true and 0 otherwise)
- ▶ Advantages:
 - ▶ Can do binary or multiclass classification.
 - ▶ Simple to compute.
 - ▶ Single value.

0/1 loss

— — —

Disadvantage: Doesn't take into account class imbalance:

- ▶ We're building a classifier to detect a rare disease.
- ▶ Assume only 1% of population is diseased.
- ▶ Diseased: $t = 1$
- ▶ Healthy: $t = 0$
- ▶ What if we always predict healthy? ($t = 0$)
- ▶ Accuracy 99%
- ▶ But classifier is rubbish!

Sensitivity and specificity

- ▶ We'll stick with our disease example.
- ▶ Need to define 4 quantities. The numbers of:
- ▶ True positives (TP) – the number of objects with $t_n^* = 1$ that are classified as $t_n = 1$ (diseased people diagnosed as diseased).
- ▶ True negatives (TN) – the number of objects with $t_n^* = 0$ that are classified as $t_n = 0$ (healthy people diagnosed as healthy).
- ▶ False positives (FP) – the number of objects with $t_n^* = 0$ that are classified as $t_n = 1$ (healthy people diagnosed as diseased).
- ▶ **False negatives (FN)** – the number of objects with $t_n^* = 1$ that are classified as $t_n = 0$ (diseased people diagnosed as healthy).

Sensitivity

— — —

$$S_e = \frac{TP}{TP + FN}$$

- ▶ The proportion of diseased people that we classify as diseased.
- ▶ The higher the better.
- ▶ In our example, $S_e = 0$.

Specificity

$$S_p = \frac{TN}{TN + FP}$$

- ▶ The proportion of healthy people that we classify as healthy.
- ▶ The higher the better.
- ▶ In our example, $S_p = 1$.

Optimising sensitivity and specificity

- ▶ We would like both to be as high as possible.
- ▶ Often increasing one will decrease the other.
- ▶ Balance will depend on application:
- ▶ e.g. diagnosis:
 - ▶ We can probably tolerate a decrease in specificity (healthy people diagnosed as diseased)....
 - ▶ ...if it gives us an increase in sensitivity (getting diseased people right).

Receiver Operating Characteristic (ROC)

— — —

- ▶ Many classification algorithms involve setting a threshold.
- ▶ e.g. Logistic Regression:

$$p(t_{new} = 1 | \mathbf{x}_{new}, \mathbf{w}) > 0.5$$


- ▶ Implies a threshold of zero (sign function)
- ▶ However, we could use any threshold we like....
- ▶ The *Receiver Operating Characteristic (ROC) curve* shows how S_e and $1 - S_p$ vary as the threshold changes.

```
In [16]: from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.datasets import load_breast_cancer

breast_cancer = load_breast_cancer()
X = breast_cancer.data
t = breast_cancer.target

X_train, X_test, y_train, y_test = train_test_split(X,t,test_size=0.30, random_state=123)
clf1 = LogisticRegression().fit(X_train, y_train)

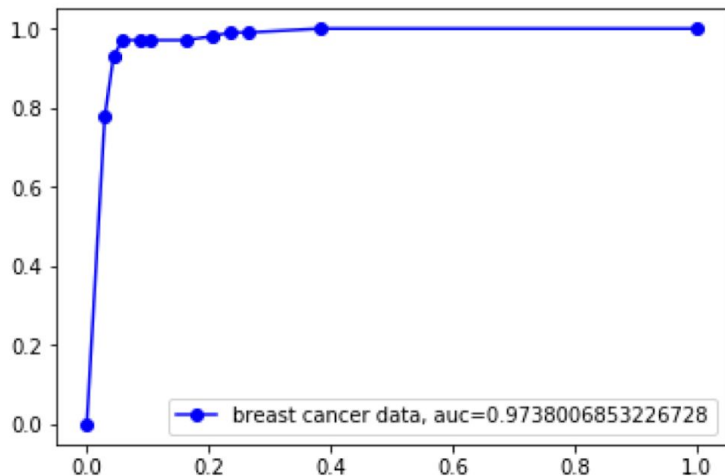
y_pred1 = clf1.predict(X_test)
y_pred_proba1 = clf1.predict_proba(X_test)[:,-1]
fpr1, tpr1, _ = metrics.roc_curve(y_test, y_pred_proba1)
auc1 = metrics.roc_auc_score(y_test, y_pred_proba1)
plt.plot(fpr1,tpr1,'ro-',label="breast cancer data, auc="+str(auc1))
plt.legend(loc=4)
plt.show()
```



Try it on a breast cancer dataset
Plot ROC of a Logistic Regression model

```
In [19]: clf2 = KNeighborsClassifier(10).fit(X_train, y_train)

y_pred2 = clf2.predict(X_test)
y_pred_proba2 = clf2.predict_proba(X_test)[: ,1]
fpr2, tpr2, _ = metrics.roc_curve(y_test, y_pred_proba2)
auc2 = metrics.roc_auc_score(y_test, y_pred_proba2)
plt.plot(fpr2,tpr2,'bo-',label="breast cancer data, auc="+str(auc2))
plt.legend(loc=4)
plt.show()
```

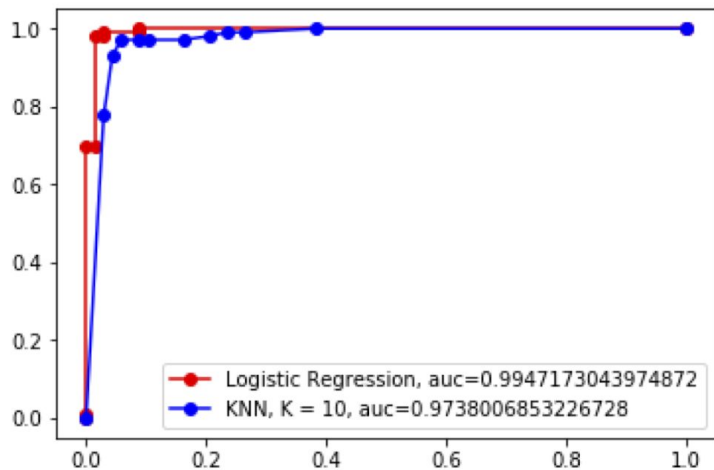


**The same data with
KNN (K =10)**

Overlay two ROC plots

— — —

```
In [21]: plt.plot(fpr1,tpr1,'ro-',label="Logistic Regression, auc="+str(auc1))  
plt.plot(fpr2,tpr2,'bo-',label="KNN, K = 10, auc="+str(auc2))  
plt.legend(loc=4)  
plt.show()
```



What is happening?

— — —

Can I have 7 volunteers?

Confusion matrix

— — —

The quantities we used to compute S_e and S_p can be neatly summarised in a table:

		True class	
		1	0
Predicted class	1	TP	FP
	0	FN	TN

- ▶ This is known as a *confusion matrix*
- ▶ It is particularly useful for multi-class classification.
- ▶ Tells us where the mistakes are being made.
- ▶ Note that normalising columns gives us S_e and S_p

Confusion matrix, example

- ▶ 20 newsgroups data.
- ▶ Thousands of documents from 20 classes (newsgroups)

			True class										
			10	11	12	13	14	15	16	18	18	19	20
Predicted class	1	...	4	2	0	2	10	4	7	1	12	7	47
	2	...	0	0	4	18	7	8	2	0	1	1	3
	3	...	0	0	1	0	1	0	1	0	0	0	0
	4	...	1	0	1	28	3	0	0	0	0	0	0
	⋮												
	16	...	3	2	2	5	17	4	376	3	7	2	68
	17	...	1	0	9	0	3	1	3	325	3	95	19
	18	...	2	1	0	2	6	2	1	2	325	4	5
	19	...	8	4	8	0	10	21	1	16	19	185	7
	20	...	0	0	1	0	1	1	2	4	0	1	92

- ▶ Algorithm is getting 'confused' between classes 20 and 16, and 19 and 17.
 - ▶ 17: talk.politics.guns
 - ▶ 19: talk.politics.misc
 - ▶ 16: talk.religion.misc
 - ▶ 20: soc.religion.christian
- ▶ Maybe these should be just one class?
- ▶ Maybe we need more data in these classes?
- ▶ Confusion matrix helps us direct our efforts to improving the classifier.